

# Amazon Order Tracking

# Agenda

- Who is this guy?
- Problem statement
- How we solved it
- Lessons

# Who am I?

- Raymond Smith (rasm@amazon.com)
- Software Developer 20+ years; 3 years a manager
- Work on Amazon's Grocery Business
- We are [looking to hire](#) 2022 graduates for 2023 start

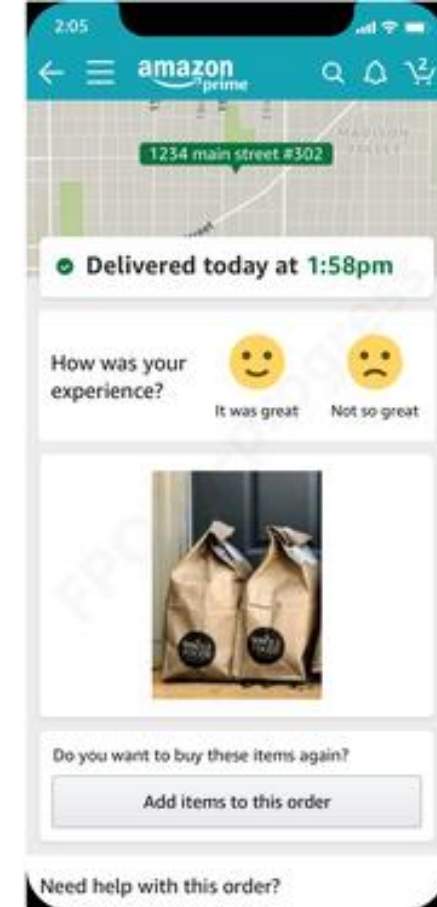
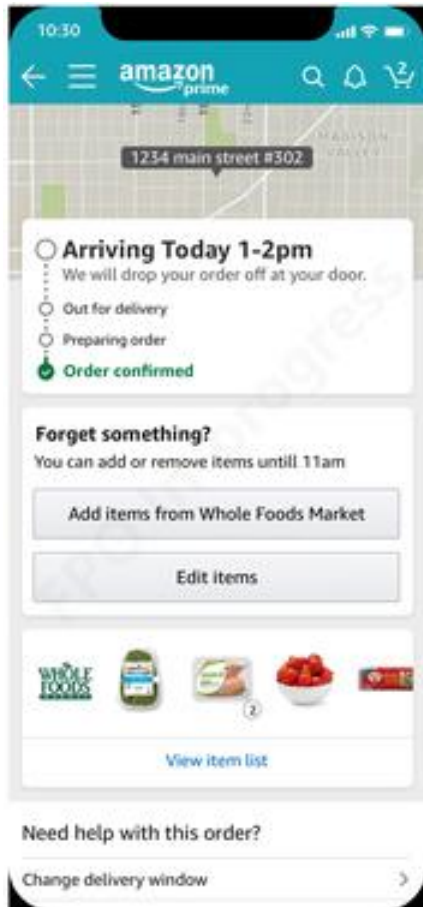
# Context: Amazon Grocery

- Amazon supports sales by its own Grocery Brands, and by third-parties using Amazon.com
- Available in: US, UK, France, Germany, Italy, Spain, Singapore, India, and Japan. More being added all the time.
- Grocery has a large development organization (more than 500 people). Even more developers in greater Amazon (tens of thousands).
- There is no way to coordinate that many developers – so we distribute the work both at a people level ([Two Pizza Teams](#)) and a Technical Level (loosely coupled, message based systems).

# Customer Problem: Where are my Groceries?

- At Amazon we always [work-backwards](#) from Customer.
- Grocery Customers have some unique needs:
  - They are relying on their order to feed themselves and their families.
  - If an order cannot be delivered, or items are missing, they likely cannot wait and will have to arrange an alternative.
- Amazon's existing notifications helped but were not great:
  - They didn't provide enough detail on where the order was at
  - They were not interactive
- Our goal was to provide a better customer experience for Amazon Customers

# Problem: Display Status of Customer's Order



# Complications

## Complications:

- 20+ services involved in fulfilling an order
- No one service owns the live state of the order
- Variations between lines of business (when events are sent)
- We don't know exactly all of the things we want to track
- Existing Order Page is on older technology

## Things in our favour:

- Loose coupling and two pizza teams let us move fast
- Rich events available with reasonably well defined formats
- Some strong common concepts: Customer ID, Order IDs, etc

# Core Technical Problem

Core problem: form a unified view by integrating events from multiple systems

Subproblems:

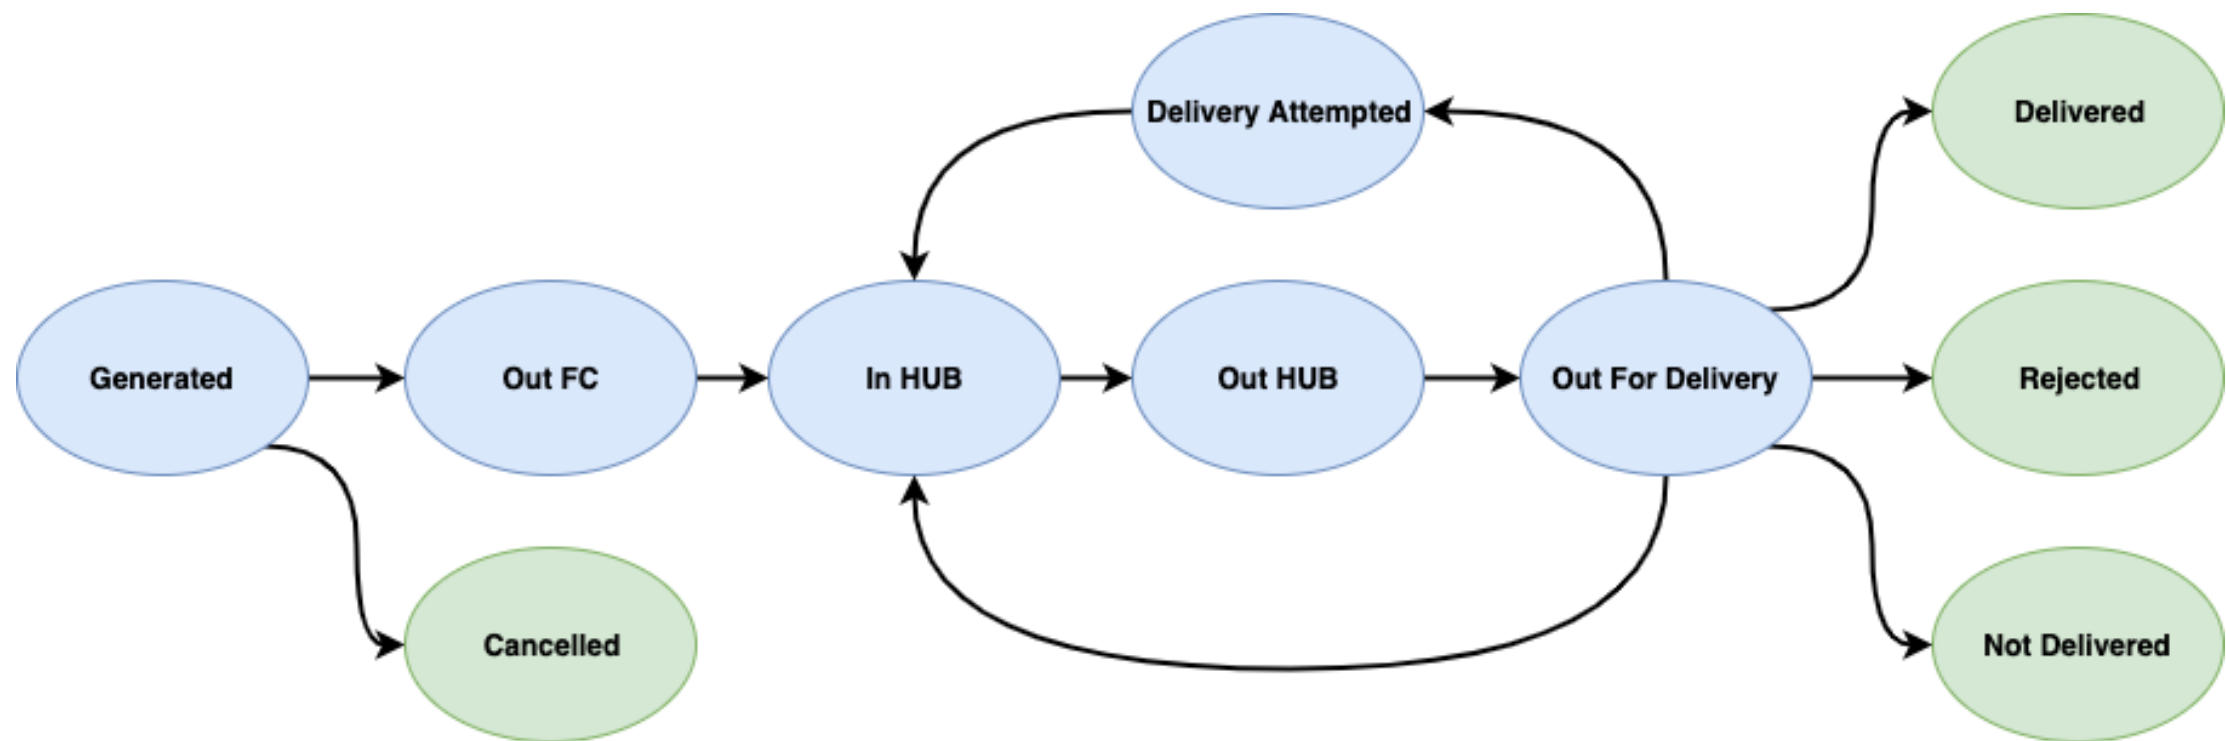
- Out of order events
- Handling missing data
- Source systems constantly evolving: new events, milestones, etc
- Keeping track of state

We already had an event bus platform for the unified view, but we did not have anything which kept track of state.

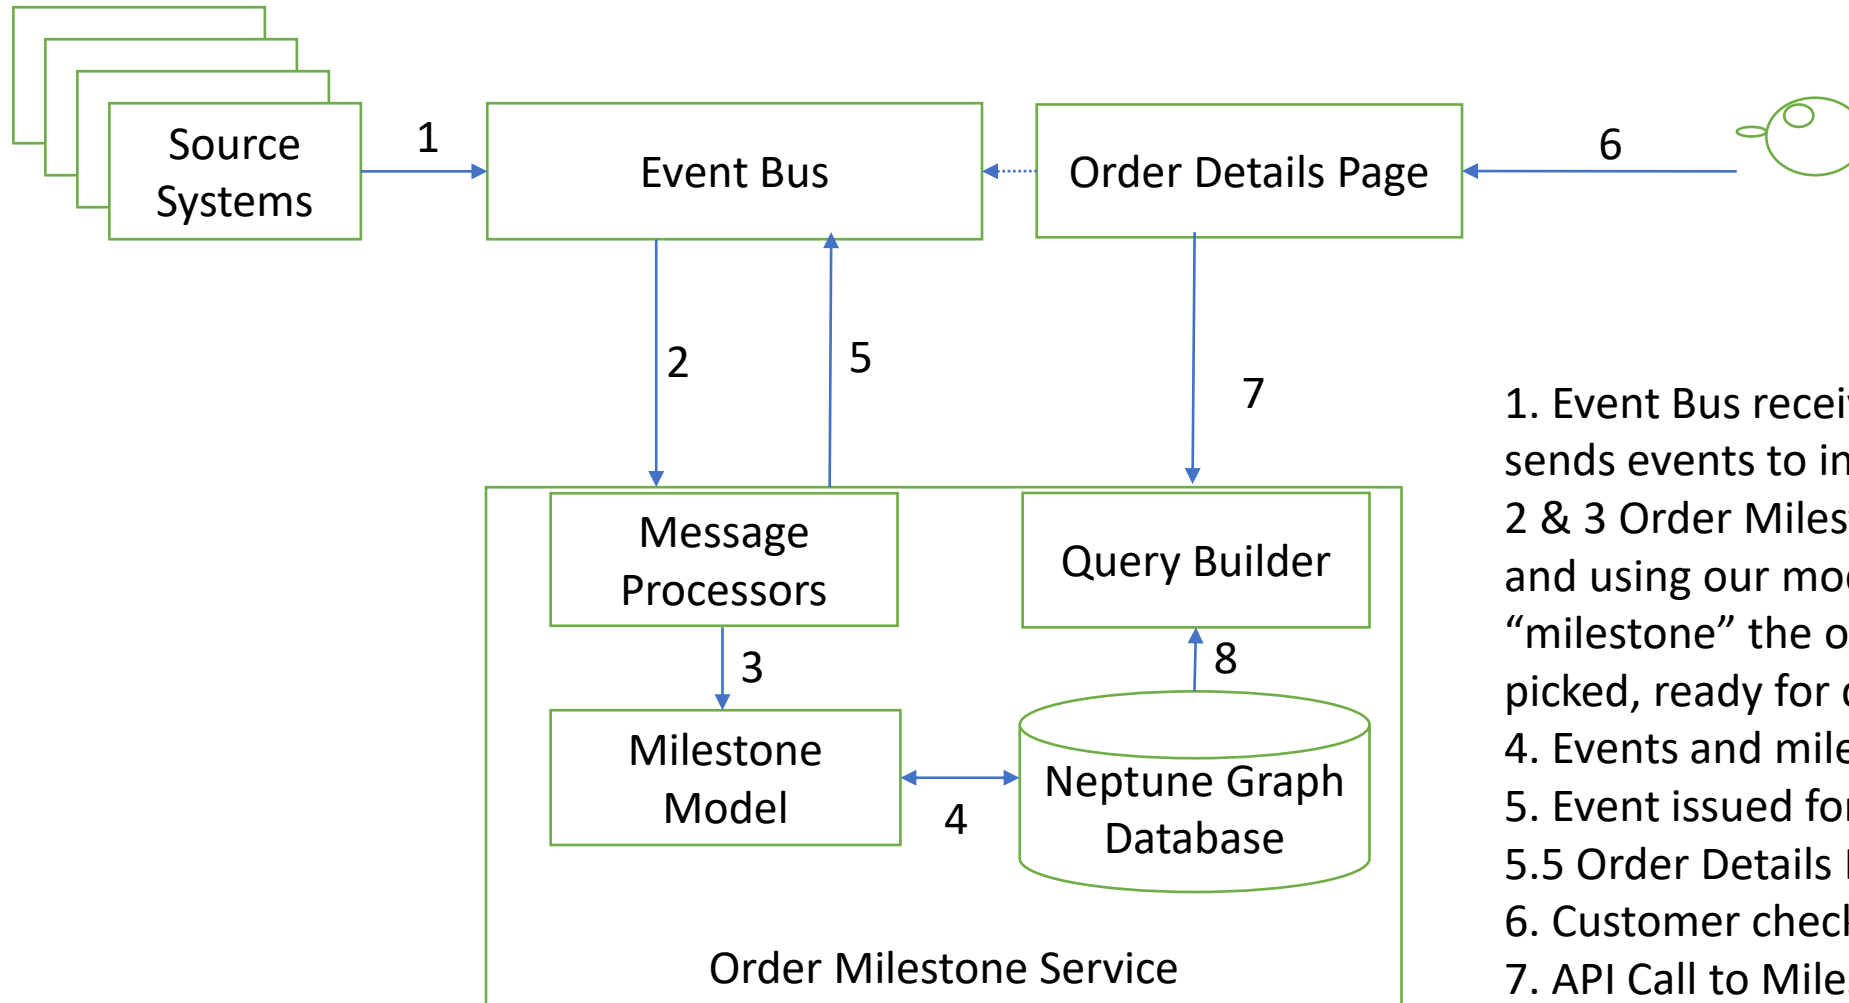


# High Level Approach

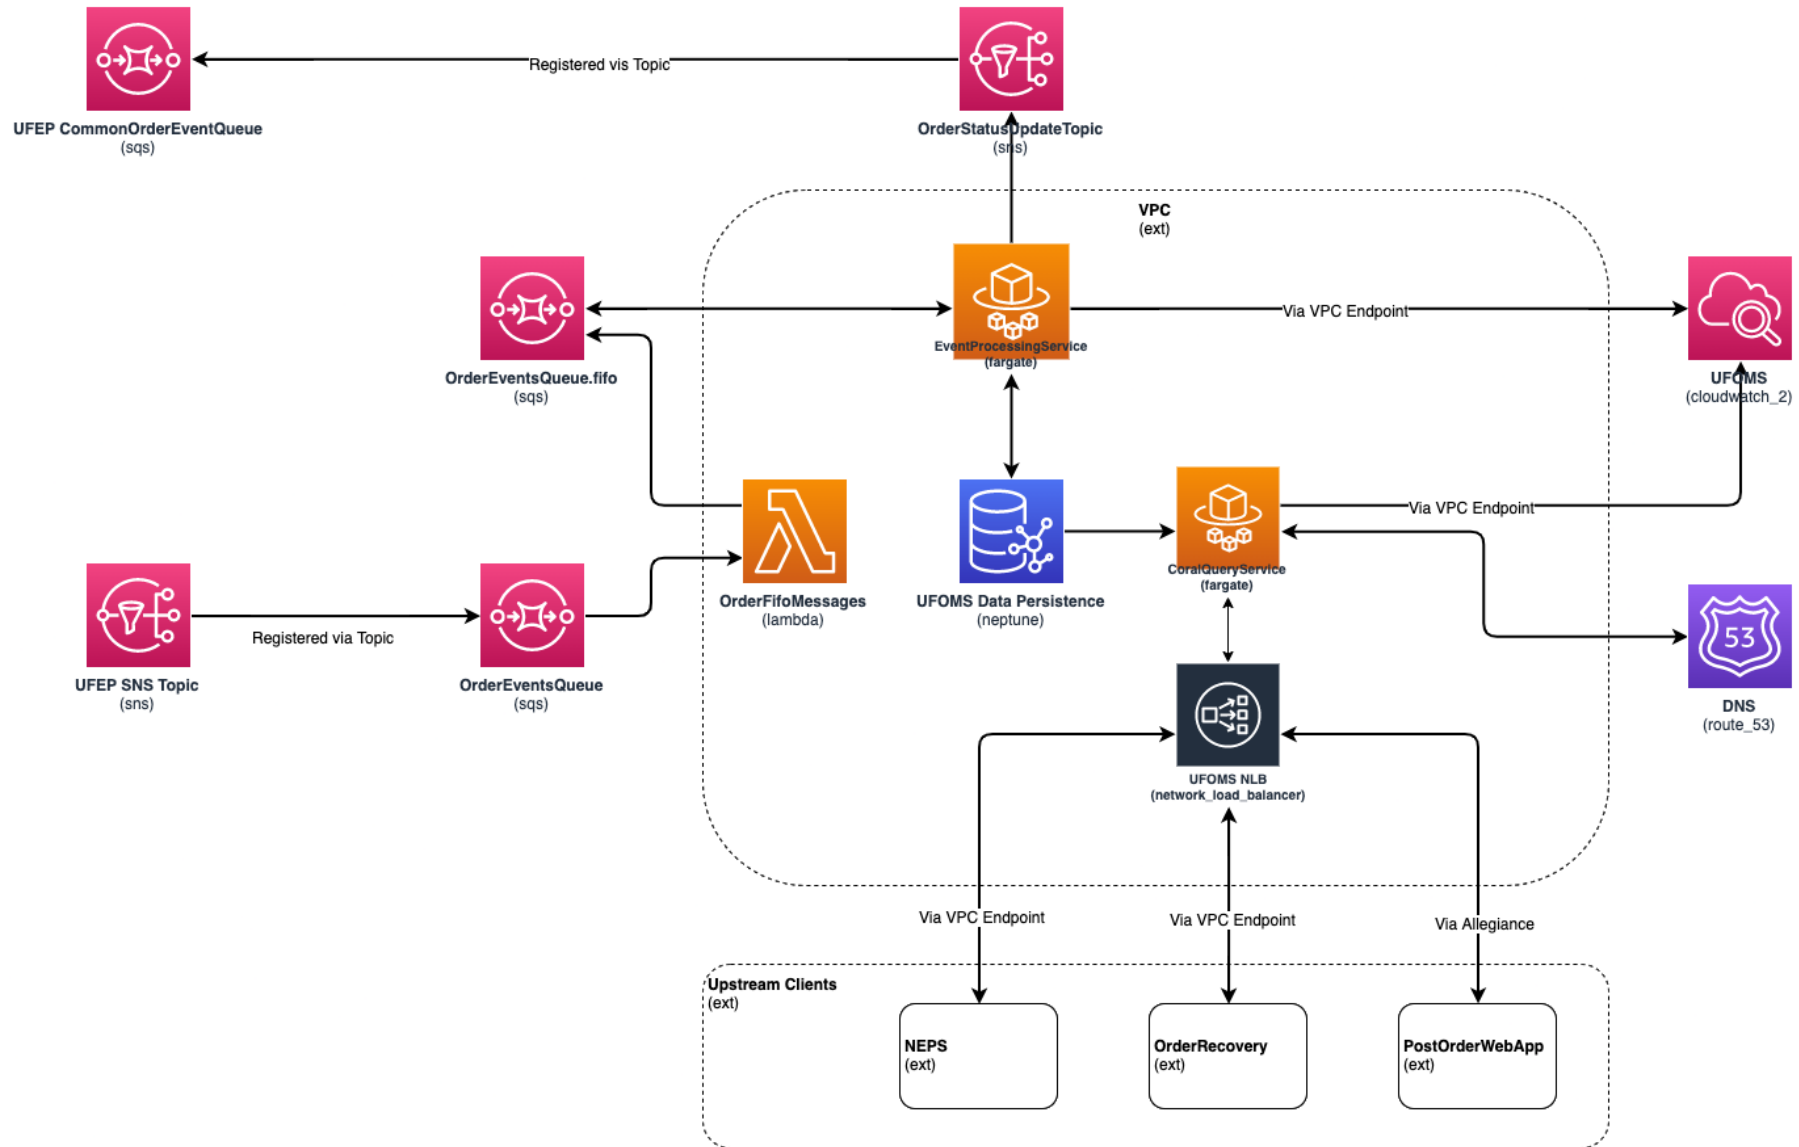
- Define a common model for the key states of an order
- Gather all the events into a common form
- Rules to trigger when different states detected
- UI to listen for state triggers for display



# (Simplified) Architecture Diagram



1. Event Bus receives events, normalizes, and sends events to interested systems.
- 2 & 3 Order Milestone Service receives events, and using our model determines which "milestone" the order has reached (created, picked, ready for delivery, etc)
4. Events and milestones stored in database
5. Event issued for new milestone
- 5.5 Order Details Page listens for milestone events
6. Customer checks on order status
7. API Call to Milestone Service
8. Queries for state in Neptune, data returned



# For Reference: Technologies we used

- AWS:
  - Compute: [Lambda](#) (serverless compute), [Fargate](#) (container compute)
  - Persistence: [Neptune](#)
  - Networking: [ALB/ELB](#), [Route 53](#), and [VPC](#)
  - Messaging: [SNS](#) and [SQS](#)
- [Apache Tinker Pop](#) & [Gremlin](#)
- Java

# Processing Events

- Distributed Systems 101: Networks are not reliable, event ordering cannot be guaranteed, the same event may be received multiple times. (This is the first [Fallacies of Distributed Computing](#).)
- Many ways to handle this
- Our approach:
  - Idempotent processing: the same event should result in the same change to the database. Achieved by using business keys (ie order id, etc).
  - Infer missing events: if we receive a “cancel order event” we can assume there was a “create order event” and fill it with placeholder information until we receive the “create event”.

# Simplified (Faked) Example

- When Events are in order:

```
// order created event
{ Id: "Order A", Ordered: [ { ASIN: "B06WW2YGSJ", QTY: 2 },
{ASIN: "B00AXYMRYA", QTY: 1} ], Status: "Created" }
// warehouse starts work event
{Id: "Order A", Ordered: [{ ASIN: "B06WW2YGSJ", QTY: 2 },
{ASIN: "B00AXYMRYA", QTY: 1}], Status: "Processing" } ],
Status: "Processing" }
// no oranges event
{"Order A", Ordered: [{ ASIN: "B06WW2YGSJ", QTY: 2 }, {ASIN:
"B00AXYMRYA", QTY: 1}], Shorted: [{ASIN: "B00AXYMRYA", QTY:
1}], Status: "Processing" }
// order is out for delivery
{ Id: "Order A", {"Order A", Ordered: [{ ASIN: "B06WW2YGSJ",
QTY: 2 }, {ASIN: "B00AXYMRYA", QTY: 1}], Shorted: [{ASIN:
"B00AXYMRYA", QTY: 1}], Status: "Out For Delivery" }
```

# Simplified (Faked) Example

- When Events are in *out* order (out for delivery comes first):

```
// order is out for delivery
{ Id: "Order A", Ordered: PLACEHOLDER, Status: "Out For Delivery" }
// order created event
{ Id: "Order A", Ordered: [ { ASIN: "B06WW2YGSJ", QTY: 2 },
{ASIN: "B00AXYMRYA", QTY: 1} ], Status: "Out For Delivery" }
// warehouse starts work event
{Id: "Order A", Ordered: [{ ASIN: "B06WW2YGSJ", QTY: 2 },
{ASIN: "B00AXYMRYA", QTY: 1}], Status: "Processing" } ],
Status: "Out For Delivery" }
// no oranges event
{Id: "Order A", Ordered: [{ ASIN: "B06WW2YGSJ", QTY: 2 },
{ASIN: "B00AXYMRYA", QTY: 1}], Shorted: [{ASIN:
"B00AXYMRYA", QTY: 1}], Status: "Out For Delivery" }
```



# Processing Events: a changing world

- Problem: the source systems and business processes are highly changeable.
- We did not want to have to rework our designs whenever there was a change. So we went looking for a flexible data-store.

# Neptune Graph Database

- Graph Database was a natural fit for a “graph of events”
- We could use Gremlin queries to determine milestones by matching sequences of events.
- Very flexible schema allowed us to add events at will.

# Neptune Controversy

- Introducing Neptune proved very controversial.
- Most Brisbane developers preferred using DynamoDB, and had gotten very good at using it for diverse purposes.
- Neptune has some serious limitations compared to DynamoDB:  
expensive, requires downtime for maintenance (up to 8 hours / year),  
requires developers to tune type of CPU and memory used.  
DynamoDB is cheap, requires no downtime, and only requires no complex tuning.





# So, why stick with Neptune?

- A graph database was a natural fit for the data we were processing.
- Innovating requires us to move outside our comfort zones.
- How has Neptune turned out?
  - AWS has been making it better and better
  - The more flexible design has paid off as it allowed us to incorporate new very easily (around 1 – 2 day, versus 2 -- 3 weeks for similar changes elsewhere).

# Aside: Radical Innovation

- Software Development is a creative activity. It is still more a craft than an engineering discipline. People become stuck in “local maximas” of technology they know well.
- Whenever you are making a radical break you need to:
  - double check that you are innovating for your customer – ensure there is genuine benefit
  - show empathy to your colleagues, actively listen to their concerns
  - have a viable plan B if your new technology does not work out
  - define success criteria, and be humble
    - if your plan does not work out, switch to plan b
  - be brave enough to be misunderstood

# The Progress Tracker Team



# Key Takeaways

- Always start with your customer, and work backwards
- Don't be afraid to try new things – but be humble and have a plan B
- Message Based Systems are powerful way to decouple systems and teams
- Distributed Systems 101: the network is not reliable – plan for that
  - Idempotence
  - Inferring state

# Discussion

Questions?

Comments?

Thoughts?



# Bonus Material: Books to Read

These three books teach timeless lessons for programmers.

